

Databricks Genie

Giving AI the context to get it right

Mir Mubasher · Data Scientist

LEVEL SET

What is Databricks?



Store



Govern



Use

Creators of

Apache Spark, Delta Lake & MLflow

20,000+

organizations worldwide

70%

of the Fortune 500

WHERE GENIE LIVES

How the Genie family fits together

Where people ask

Genie One web | Slack | Microsoft Teams | iOS & Android | MCP App

Genie One

The front door: ask, discover, act. Looks for a matching agent first.

No match? Searches permitted dashboards, queries & metric views

routes to

Genie Agents

the trusted domain specialists

- **Sales agent** Curated tables & metric views
- **Support agent** Trusted SQL & benchmarks

Genie Code

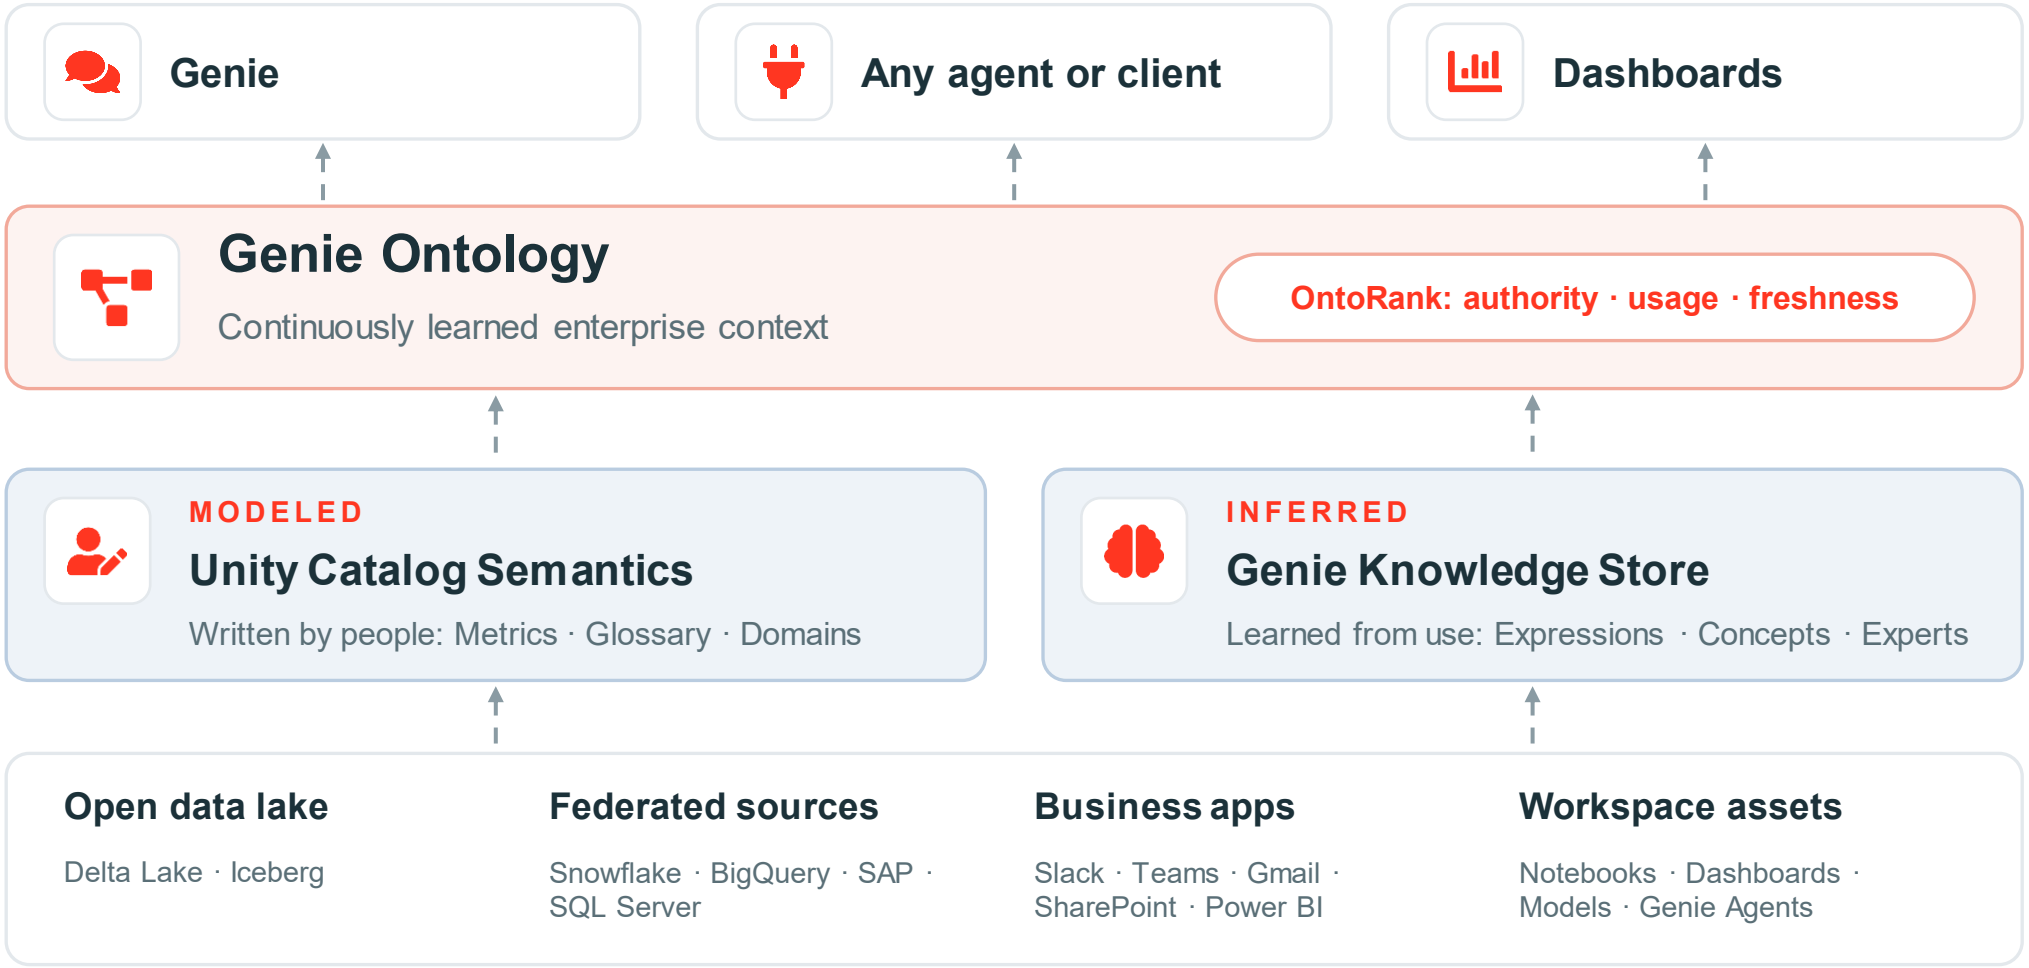
The builder: creates and tunes agents, SQL & pipelines

Genie Ontology The brain underneath: a living graph of metrics, terms & trusted sources, ranked by authority

Unity Catalog Permissions enforced per user on every answer

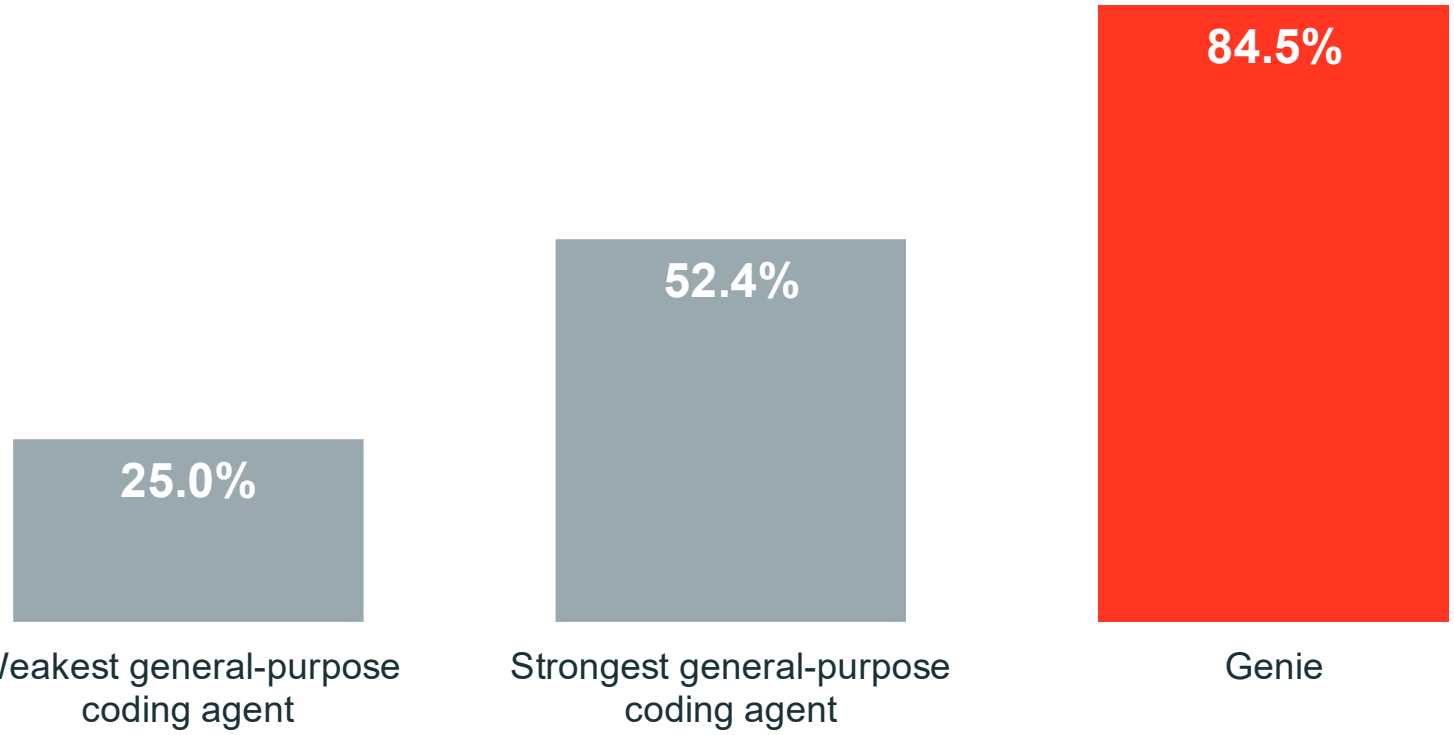
HOW GENIE THINKS

Genie Ontology: one context layer for everything



DOES IT WORK?

Context beats raw model power



+32 pts

more accurate than the strongest coding agent

~2×

faster than the strongest coding agent

Source: Databricks internal benchmark · 28 real-world data-analysis questions · June 2026 · competitors anonymized

HOW GENIE THINKS

From question to trusted answer in five steps

01

Context selection

Picks the relevant tables, metrics, example SQL & instructions

02

Value matching

Maps your words to stored values, like "Florida" to "FL"

03

SQL generation

Writes read-only SQL, or asks a clarifying question if unsure

04

Execution

Runs on the SQL warehouse; access checked as the end user

05

Response

Explanation, the SQL, a results table & an optional chart

OntoRank

decides which sources are most trustworthy

Unity Catalog

row filters & column masks apply per person

“Unity Catalog decides what you can see. OntoRank decides what's trustworthy.”

WHERE ACCURACY COMES FROM

Errors flow left to right

1 · Data model → 2 · Governance → 3 · Semantics → 4 · Prompt

LAYER	FAILURE WHEN WEAK	PLATFORM FIX
1 · Data model	Fan-out joins, double-counted totals	Denormalized gold views; a small set of curated tables
2 · Governance	PII exposure, or SQL failing on missing USE CATALOG / USE SCHEMA grants	UC row/column security, ABAC, explicit group grants
3 · Semantics	Metric drift: two users, two "revenue" numbers	Metric views, primary/foreign keys in UC, trusted queries
4 · Prompt	Acronyms and jargon misread	Agent instructions, synonyms, example SQL

A prompt fix can't repair a problem further upstream.

RECAP

Context. Control. Cost. Choice.

Context

Genie Ontology: modeled + learned business meaning

Control

Unity Catalog permissions at every layer

Cost

No seat licenses; human usage free through Jan 2027

Choice

Any agent or client, open formats, any cloud

THANK YOU

Questions?

Mir Mubasher · Data Scientist

LET'S CONNECT

Scan to connect with Mir Mubasher